

# analysis\_chicago\_all

April 23, 2023

## 1 Analysis Template

### 1.1 Preprocess

```
[ ]: # resolve dependency
     # !pip install pmdarima
```

```
[ ]: import numpy as np
     import pandas as pd
     import matplotlib.pyplot as plt
     from statsmodels.tsa.stattools import adfuller
     from pandas.plotting import autocorrelation_plot
     from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
     import statsmodels.api as sm
     from pmdarima.arima import ADFTest, auto_arima
     %matplotlib inline
```

```
[ ]: data_path = "../data/crime_occurrence_per_day.csv"
     crime = "all"
     target = "Count"
     date = "Date"
     city = "chicago"
     fig_size = (20,5)
```

```
[ ]: df_by_day = pd.read_csv(data_path)
     df_by_day[date] = pd.to_datetime(df_by_day[date])
     df_by_day.set_index(date, inplace=True)
```

### 1.2 Profiling

#### 1.2.1 By day

```
[ ]: df_by_day.head()
```

```
[ ]:
      Count
Date
2001-01-01  1825
2001-01-02  1143
```

```

2001-01-03    1151
2001-01-04    1166
2001-01-05    1267

```


```
[ ]: df_by_day.describe()
```

```

[ ]:
      Count
count  8132.000000
mean   956.034801
std    285.218720
min     320.000000
25%    717.000000
50%    913.000000
75%   1207.000000
max   2033.000000

```

```
[ ]: df_by_day.plot(figsize=fig_size, title="Daily Average Occurrence " + crime)
plt.show()
```



```
[ ]: df_by_day[target].sort_values(ascending=False).head()
```

```


[ ]: Date
2003-01-01    2033
2005-01-01    1977
2007-01-01    1961
2006-01-01    1918
2020-05-31    1899
Name: Count, dtype: int64

```

### 1.2.2 By week

```
[ ]: df_by_week = pd.DataFrame(df_by_day[target].resample('W').mean())
```


```
[ ]: df_by_week.plot(
    figsize=fig_size,
    title="Weekly Average Occurrence " + crime)
plt.show()
```



### 1.2.3 By month

```
[ ]: df_by_month = pd.DataFrame(df_by_day[target].resample('M').mean())
```

```
[ ]: df_by_month.plot(
    figsize=fig_size,
    title="Monthly Average Occurrence " + crime)
plt.show()
```



## 1.3 Analysis

```
[ ]: #Ho: It is non stationary
    #H1: It is stationary

def adfuller_test(count):
    result=adfuller(count)
```

```

labels = ['ADF Test Statistic','p-value','#Lags Used','Number of_
↳Observations Used']
for value,label in zip(result,labels):
    print(label+' : '+str(value) )
if result[1] <= 0.05:
    print("strong evidence against the null hypothesis(Ho), reject the null_
↳hypothesis. Data has no unit root and is stationary")
else:
    print("weak evidence against null hypothesis, time series has a unit_
↳root, indicating it is non-stationary ")

```

### 1.3.1 Checking stationary

```
[ ]: adfuller_test(df_by_month[target])
```

```

ADF Test Statistic : -1.3090286063052534
p-value : 0.6250316590534809
#Lags Used : 15
Number of Observations Used : 252
weak evidence against null hypothesis, time series has a unit root, indicating
it is non-stationary

```

### 1.3.2 Checking seasonality

```
[ ]: df_by_month['seasonal_first_difference'] = df_by_month[target] -_
↳df_by_month[target].shift(12)
```

```
[ ]: adfuller_test(df_by_month['seasonal_first_difference'].dropna())
```


```

ADF Test Statistic : -2.912294910803625
p-value : 0.04393683367922969
#Lags Used : 12
Number of Observations Used : 243
strong evidence against the null hypothesis(Ho), reject the null hypothesis.
Data has no unit root and is stationary

```

```
[ ]: df_by_month['seasonal_first_difference'].plot(figsize=fig_size, title='Seasonal_
↳difference plot - ' + crime)
```


```
[ ]: <Axes: title={'center': 'Seasonal difference plot - all'}, xlabel='Date'>
```



### 1.3.3 Auto Regressive Model

```
[ ]: fig = plt.figure(figsize=(12,8))
ax1 = fig.add_subplot(211)
fig = sm.graphics.tsa.plot_acf(df_by_month['seasonal_first_difference'].iloc[13:
↪],lags=80,ax=ax1)
ax2 = fig.add_subplot(212)
fig = sm.graphics.tsa.plot_pacf(df_by_month['seasonal_first_difference'].
↪iloc[13:],lags=80,ax=ax2)
```

/Users/xuyanchong/opt/anaconda3/lib/python3.9/site-packages/statsmodels/graphics/tsaplots.py:348: FutureWarning: The default method 'yw' can produce PACF values outside of the [-1,1] interval. After 0.13, the default will change to unadjusted Yule-Walker ('ywm'). You can use this method now by setting method='ywm'.  
 warnings.warn(




### 1.3.4 Implementing Seasonal Arima Model

```
[ ]: adf_test=ADFTTest(alpha=0.05)
      adf_test.should_diff(df_by_month[target])
```

```
[ ]: (0.01, False)
```

```
[ ]: start=int(df_by_month.shape[0]*0.8)
      train=df_by_month[:start]
      test=df_by_month[start:]
      plt.figure(figsize=fig_size)
      plt.plot(train[target])
      plt.plot(test[target])
      plt.title('Training and Testing split - '+ crime)
      plt.show()
```



```
[ ]: model=auto_arima(train[target],start_p=0,d=1,start_q=0,
    max_p=10,max_d=10,max_q=10, start_P=0,
    D=1, start_Q=0, max_P=10,max_D=10,
    max_Q=10, m=12, seasonal=True,
    error_action='warn',trace=True,
    supress_warnings=True,stepwise=True,
    random_state=20,n_fits=50)
```

Performing stepwise search to minimize aic

ARIMA(0,1,0)(0,1,0)[12]	: AIC=2112.466, Time=0.04 sec
ARIMA(1,1,0)(1,1,0)[12]	: AIC=2045.289, Time=0.49 sec
ARIMA(0,1,1)(0,1,1)[12]	: AIC=1983.008, Time=3.07 sec
ARIMA(0,1,1)(0,1,0)[12]	: AIC=2048.790, Time=0.19 sec
ARIMA(0,1,1)(1,1,1)[12]	: AIC=1984.906, Time=3.66 sec
ARIMA(0,1,1)(0,1,2)[12]	: AIC=1984.841, Time=26.05 sec
ARIMA(0,1,1)(1,1,0)[12]	: AIC=2017.890, Time=0.75 sec
ARIMA(0,1,1)(1,1,2)[12]	: AIC=inf, Time=37.85 sec
ARIMA(0,1,0)(0,1,1)[12]	: AIC=2028.476, Time=1.61 sec
ARIMA(1,1,1)(0,1,1)[12]	: AIC=1979.815, Time=3.28 sec
ARIMA(1,1,1)(0,1,0)[12]	: AIC=2046.407, Time=0.47 sec
ARIMA(1,1,1)(1,1,1)[12]	: AIC=1981.654, Time=5.32 sec
ARIMA(1,1,1)(0,1,2)[12]	: AIC=1981.561, Time=27.43 sec
ARIMA(1,1,1)(1,1,0)[12]	: AIC=2015.549, Time=1.34 sec
ARIMA(1,1,1)(1,1,2)[12]	: AIC=inf, Time=55.11 sec
ARIMA(1,1,0)(0,1,1)[12]	: AIC=2004.148, Time=2.41 sec
ARIMA(2,1,1)(0,1,1)[12]	: AIC=1981.810, Time=3.78 sec
ARIMA(1,1,2)(0,1,1)[12]	: AIC=1981.812, Time=3.06 sec
ARIMA(0,1,2)(0,1,1)[12]	: AIC=1980.112, Time=1.82 sec
ARIMA(2,1,0)(0,1,1)[12]	: AIC=1996.287, Time=1.70 sec
ARIMA(2,1,2)(0,1,1)[12]	: AIC=1983.547, Time=2.83 sec
ARIMA(1,1,1)(0,1,1)[12] intercept	: AIC=1981.803, Time=1.64 sec

Best model: ARIMA(1,1,1)(0,1,1)[12]

Total fit time: 183.924 seconds

```
[ ]: model.summary()
```

```
[ ]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```


                                SARIMAX Results
=====
Dep. Variable:                  y    No. Observations:
214
Model:              SARIMAX(1, 1, 1)x(0, 1, 1, 12)    Log Likelihood
-985.908
Date:                  Sun, 23 Apr 2023    AIC
1979.815
Time:                  01:34:43    BIC
1993.029
Sample:                01-31-2001    HQIC
1985.162
                        - 10-31-2018
Covariance Type:                opg
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1           0.2290       0.116       1.967     0.049       0.001       0.457
ma.L1          -0.7716       0.075     -10.324     0.000      -0.918      -0.625
ma.S.L12       -0.6947       0.055     -12.529     0.000      -0.803      -0.586
sigma2        1021.9363     82.140      12.441     0.000     860.945     1182.927
=====
===
Ljung-Box (L1) (Q):                0.00    Jarque-Bera (JB):
14.24
Prob(Q):                0.99    Prob(JB):
0.00
Heteroskedasticity (H):            0.62    Skew:
0.15
Prob(H) (two-sided):            0.05    Kurtosis:
4.27
=====
===

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-
step).
      """
```

```
[ ]: prediction = pd.DataFrame(model.predict(n_periods = train.shape[0]),index=test.
      ↪index)
      prediction.columns = ['predicted_crime']
```



```
plt.figure(figsize=fig_size)
plt.plot(train[target],label="Training")
plt.plot(test[target],label="Test")
plt.plot(prediction,label="Predicted")
plt.legend(loc = 'upper left')
plt.savefig('../output/%s_%s_pred.jpg' % (city,crime))
plt.title('Model output - '+crime)
plt.show()
```



```
[ ]: np.sqrt(np.square(np.subtract(test[target].values,prediction['predicted_crime'].
↪values)).mean())
```

```
[ ]: 89.9247972781049
```