

analysis_chicago_type1

April 23, 2023

1 Analysis Template

1.1 Preprocess

```
[ ]: # resolve dependency
     # !pip install pmdarima
```

```
[ ]: import numpy as np
     import pandas as pd
     import matplotlib.pyplot as plt
     from statsmodels.tsa.stattools import adfuller
     from pandas.plotting import autocorrelation_plot
     from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
     import statsmodels.api as sm
     from pmdarima.arima import ADFTest, auto_arima
     %matplotlib inline
```

```
[ ]: data_path = "../data/battery_occurrence_per_day.csv"
     crime = "type1"
     target = "Count"
     date = "Date"
     city = "chicago"
     fig_size = (20,5)
```

```
[ ]: df_by_day = pd.read_csv(data_path)
     df_by_day[date] = pd.to_datetime(df_by_day[date])
     df_by_day.set_index(date, inplace=True)
```

1.2 Profiling

1.2.1 By day

```
[ ]: df_by_day.head()
```

```
[ ]:
```


| Date | Count |
|------------|-------|
| 2001-01-01 | 296 |
| 2001-01-02 | 143 |

```
2001-01-03    165
2001-01-04    173
2001-01-05    178
```

```
[ ]: df_by_day.describe()
```

```
[ ]:
      Count
count  8132.000000
mean   174.757993
std     61.009520
min     53.000000
25%    127.000000
50%    163.000000
75%    216.000000
max     423.000000
```

```
[ ]: df_by_day.plot(figsize=fig_size, title="Daily Average Occurrence " + crime)
plt.show()
```




```
[ ]: df_by_day[target].sort_values(ascending=False).head()
```

```
[ ]: Date
2001-06-17    423
2002-06-23    411
2002-07-04    409
2002-06-01    397
2002-07-20    393
Name: Count, dtype: int64
```

1.2.2 By week

```
[ ]: df_by_week = pd.DataFrame(df_by_day[target].resample('W').mean())
```


```
[ ]: df_by_week.plot(
    figsize=fig_size,
    title="Weekly Average Occurrence " + crime)
plt.show()
```



1.2.3 By month

```
[ ]: df_by_month = pd.DataFrame(df_by_day[target].resample('M').mean())
```

```
[ ]: df_by_month.plot(
    figsize=fig_size,
    title="Monthly Average Occurrence " + crime)
plt.show()
```



1.3 Analysis

```
[ ]: #Ho: It is non stationary
    #H1: It is stationary

def adfuller_test(count):
    result=adfuller(count)
```

```

labels = ['ADF Test Statistic','p-value','#Lags Used','Number of_
↳Observations Used']
for value,label in zip(result,labels):
    print(label+' : '+str(value) )
if result[1] <= 0.05:
    print("strong evidence against the null hypothesis(Ho), reject the null_
↳hypothesis. Data has no unit root and is stationary")
else:
    print("weak evidence against null hypothesis, time series has a unit_
↳root, indicating it is non-stationary ")

```

1.3.1 Checking stationary

```
[ ]: adfuller_test(df_by_month[target])
```

```

ADF Test Statistic : -1.1257894251667415
p-value : 0.7046785869157075
#Lags Used : 14
Number of Observations Used : 253
weak evidence against null hypothesis, time series has a unit root, indicating
it is non-stationary

```

1.3.2 Checking seasonality

```
[ ]: df_by_month['seasonal_first_difference'] = df_by_month[target] -_
↳df_by_month[target].shift(12)
```

```
[ ]: adfuller_test(df_by_month['seasonal_first_difference'].dropna())
```


```

ADF Test Statistic : -4.335873703851938
p-value : 0.0003847880016246871
#Lags Used : 12
Number of Observations Used : 243
strong evidence against the null hypothesis(Ho), reject the null hypothesis.
Data has no unit root and is stationary

```

```
[ ]: df_by_month['seasonal_first_difference'].plot(figsize=fig_size, title='Seasonal_
↳difference plot - ' + crime)
```

```
[ ]: <Axes: title={'center': 'Seasonal difference plot - type1'}, xlabel='Date'>
```




1.3.3 Auto Regressive Model

```
[ ]: fig = plt.figure(figsize=(12,8))
      ax1 = fig.add_subplot(211)
      fig = sm.graphics.tsa.plot_acf(df_by_month['seasonal_first_difference'].iloc[13:
      ↪],lags=80,ax=ax1)
      ax2 = fig.add_subplot(212)
      fig = sm.graphics.tsa.plot_pacf(df_by_month['seasonal_first_difference'].
      ↪iloc[13:],lags=80,ax=ax2)
```

/Users/xuyanchong/opt/anaconda3/lib/python3.9/site-packages/statsmodels/graphics/tsaplots.py:348: FutureWarning: The default method 'yw' can produce PACF values outside of the [-1,1] interval. After 0.13, the default will change to unadjusted Yule-Walker ('ywm'). You can use this method now by setting method='ywm'.

```
warnings.warn(
```




1.3.4 Implementing Seasonal Arima Model

```
[ ]: adf_test=ADFTTest(alpha=0.05)
      adf_test.should_diff(df_by_month[target])
```

```
[ ]: (0.01, False)
```

```
[ ]: start=int(df_by_month.shape[0]*0.8)
      train=df_by_month[:start]
      test=df_by_month[start:]
      plt.figure(figsize=fig_size)
      plt.plot(train[target])
      plt.plot(test[target])
      plt.title('Training and Testing split - '+ crime)
      plt.show()
```



```
[ ]: model=auto_arima(train[target],start_p=0,d=1,start_q=0,
    max_p=10,max_d=10,max_q=10, start_P=0,
    D=1, start_Q=0, max_P=10,max_D=10,
    max_Q=10, m=12, seasonal=True,
    error_action='warn',trace=True,
    supress_warnings=True,stepwise=True,
    random_state=20,n_fits=50)
```

Performing stepwise search to minimize aic

| | |
|-----------------------------------|--------------------------------|
| ARIMA(0,1,0)(0,1,0)[12] | : AIC=1599.851, Time=0.07 sec |
| ARIMA(1,1,0)(1,1,0)[12] | : AIC=1521.442, Time=0.81 sec |
| ARIMA(0,1,1)(0,1,1)[12] | : AIC=1468.648, Time=3.16 sec |
| ARIMA(0,1,1)(0,1,0)[12] | : AIC=1533.339, Time=0.17 sec |
| ARIMA(0,1,1)(1,1,1)[12] | : AIC=1470.446, Time=3.66 sec |
| ARIMA(0,1,1)(0,1,2)[12] | : AIC=1470.306, Time=22.09 sec |
| ARIMA(0,1,1)(1,1,0)[12] | : AIC=1494.488, Time=0.93 sec |
| ARIMA(0,1,1)(1,1,2)[12] | : AIC=inf, Time=28.72 sec |
| ARIMA(0,1,0)(0,1,1)[12] | : AIC=1518.288, Time=1.57 sec |
| ARIMA(1,1,1)(0,1,1)[12] | : AIC=1464.449, Time=3.33 sec |
| ARIMA(1,1,1)(0,1,0)[12] | : AIC=1531.298, Time=0.54 sec |
| ARIMA(1,1,1)(1,1,1)[12] | : AIC=1466.355, Time=3.95 sec |
| ARIMA(1,1,1)(0,1,2)[12] | : AIC=1466.295, Time=25.90 sec |
| ARIMA(1,1,1)(1,1,0)[12] | : AIC=1493.450, Time=1.47 sec |
| ARIMA(1,1,1)(1,1,2)[12] | : AIC=inf, Time=36.14 sec |
| ARIMA(1,1,0)(0,1,1)[12] | : AIC=1487.935, Time=2.43 sec |
| ARIMA(2,1,1)(0,1,1)[12] | : AIC=1464.919, Time=4.66 sec |
| ARIMA(1,1,2)(0,1,1)[12] | : AIC=1465.020, Time=4.92 sec |
| ARIMA(0,1,2)(0,1,1)[12] | : AIC=1466.224, Time=3.90 sec |
| ARIMA(2,1,0)(0,1,1)[12] | : AIC=1480.192, Time=3.27 sec |
| ARIMA(2,1,2)(0,1,1)[12] | : AIC=1466.866, Time=4.89 sec |
| ARIMA(1,1,1)(0,1,1)[12] intercept | : AIC=1466.165, Time=3.09 sec |

Best model: ARIMA(1,1,1)(0,1,1)[12]

Total fit time: 159.672 seconds

```
[ ]: model.summary()
```

```
[ ]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                                SARIMAX Results
=====
Dep. Variable:                  y      No. Observations:
214
Model:              SARIMAX(1, 1, 1)x(0, 1, 1, 12)      Log Likelihood
-728.224
Date:                  Sun, 23 Apr 2023      AIC
1464.449
Time:                  01:34:17      BIC
1477.662
Sample:                01-31-2001      HQIC
1469.795
                        - 10-31-2018
Covariance Type:                opg
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1          0.2699      0.097       2.788      0.005      0.080      0.460
ma.L1         -0.8472      0.052     -16.196      0.000     -0.950     -0.745
ma.S.L12      -0.6542      0.071      -9.177      0.000     -0.794     -0.515
sigma2        79.0241      8.944       8.835      0.000     61.494     96.554
=====
===
Ljung-Box (L1) (Q):                0.09      Jarque-Bera (JB):
0.97
Prob(Q):                0.77      Prob(JB):
0.62
Heteroskedasticity (H):            0.41      Skew:
-0.01
Prob(H) (two-sided):            0.00      Kurtosis:
2.66
=====
===
```


Warnings:

```
[1] Covariance matrix calculated using the outer product of gradients (complex-
step).
```

```
[ ]: prediction = pd.DataFrame(model.predict(n_periods = train.shape[0]),index=test.
      ↪index)
      prediction.columns = ['predicted_crime']
```



```
plt.figure(figsize=fig_size)
plt.plot(train[target],label="Training")
plt.plot(test[target],label="Test")
plt.plot(prediction,label="Predicted")
plt.legend(loc = 'upper left')
plt.savefig('../output/%s_%s_pred.jpg' % (city,crime))
plt.title('Model output - '+crime)
plt.show()
```



```
[ ]: np.sqrt(np.square(np.subtract(test[target].values,prediction['predicted_crime'].
↪values)).mean())
```

```
[ ]: 17.381555810118396
```