

analysis_chicago_type2

April 23, 2023

1 Analysis Template

1.1 Preprocess

```
[ ]: # resolve dependency
     # !pip install pmdarima
```

```
[ ]: import numpy as np
     import pandas as pd
     import matplotlib.pyplot as plt
     from statsmodels.tsa.stattools import adfuller
     from pandas.plotting import autocorrelation_plot
     from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
     import statsmodels.api as sm
     from pmdarima.arima import ADFTest, auto_arima
     %matplotlib inline
```

```
[ ]: data_path = "../data/theft_occurrence_per_day.csv"
     crime = "type2"
     target = "Count"
     date = "Date"
     city = "chicago"
     fig_size = (20,5)
```

```
[ ]: df_by_day = pd.read_csv(data_path)
     df_by_day[date] = pd.to_datetime(df_by_day[date])
     df_by_day.set_index(date, inplace=True)
```

1.2 Profiling

1.2.1 By day

```
[ ]: df_by_day.head()
```

```
[ ]:
```


Date	Count
2001-01-01	412
2001-01-02	221

```
2001-01-03    226
2001-01-04    243
2001-01-05    265
```

```
[ ]: df_by_day.describe()
```

```
[ ]:
      Count
count  8132.000000
mean   201.678554
std     59.238356
min     38.000000
25%    160.000000
50%    199.000000
75%    241.000000
max     567.000000
```

```
[ ]: df_by_day.plot(figsize=fig_size, title="Daily Average Occurrence " + crime)
plt.show()
```




```
[ ]: df_by_day[target].sort_values(ascending=False).head()
```

```
[ ]: Date
2008-01-01    567
2007-01-01    546
2006-01-01    520
2003-01-01    472
2005-01-01    466
Name: Count, dtype: int64
```

1.2.2 By week

```
[ ]: df_by_week = pd.DataFrame(df_by_day[target].resample('W').mean())
```


```
[ ]: df_by_week.plot(
    figsize=fig_size,
    title="Weekly Average Occurrence " + crime)
plt.show()
```



1.2.3 By month

```
[ ]: df_by_month = pd.DataFrame(df_by_day[target].resample('M').mean())
```

```
[ ]: df_by_month.plot(
    figsize=fig_size,
    title="Monthly Average Occurrence " + crime)
plt.show()
```



1.3 Analysis

```
[ ]: #Ho: It is non stationary
    #H1: It is stationary

def adfuller_test(count):
    result=adfuller(count)
```

```

labels = ['ADF Test Statistic','p-value','#Lags Used','Number of_
↳Observations Used']
for value,label in zip(result,labels):
    print(label+' : '+str(value) )
if result[1] <= 0.05:
    print("strong evidence against the null hypothesis(Ho), reject the null_
↳hypothesis. Data has no unit root and is stationary")
else:
    print("weak evidence against null hypothesis, time series has a unit_
↳root, indicating it is non-stationary ")

```

1.3.1 Checking stationary

```
[ ]: adfuller_test(df_by_month[target])
```

```

ADF Test Statistic : -1.2545143770290907
p-value : 0.6497168328079318
#Lags Used : 13
Number of Observations Used : 254
weak evidence against null hypothesis, time series has a unit root, indicating
it is non-stationary

```

1.3.2 Checking seasonality

```
[ ]: df_by_month['seasonal_first_difference'] = df_by_month[target] -_
↳df_by_month[target].shift(12)
```

```
[ ]: adfuller_test(df_by_month['seasonal_first_difference'].dropna())
```


```

ADF Test Statistic : -4.324409680804535
p-value : 0.0004029082855870488
#Lags Used : 12
Number of Observations Used : 243
strong evidence against the null hypothesis(Ho), reject the null hypothesis.
Data has no unit root and is stationary

```

```
[ ]: df_by_month['seasonal_first_difference'].plot(figsize=fig_size, title='Seasonal_
↳difference plot - ' + crime)
```

```
[ ]: <Axes: title={'center': 'Seasonal difference plot - type2'}, xlabel='Date'>
```




1.3.3 Auto Regressive Model

```
[ ]: fig = plt.figure(figsize=(12,8))
ax1 = fig.add_subplot(211)
fig = sm.graphics.tsa.plot_acf(df_by_month['seasonal_first_difference'].iloc[13:
↪],lags=80,ax=ax1)
ax2 = fig.add_subplot(212)
fig = sm.graphics.tsa.plot_pacf(df_by_month['seasonal_first_difference'].
↪iloc[13:],lags=80,ax=ax2)
```

/Users/xuyanchong/opt/anaconda3/lib/python3.9/site-packages/statsmodels/graphics/tsaplots.py:348: FutureWarning: The default method 'yw' can produce PACF values outside of the $[-1,1]$ interval. After 0.13, the default will change to unadjusted Yule-Walker ('ywm'). You can use this method now by setting method='ywm'.

```
warnings.warn(
```




1.3.4 Implementing Seasonal Arima Model

```
[ ]: adf_test=ADFTTest(alpha=0.05)
      adf_test.should_diff(df_by_month[target])
```

```
[ ]: (0.01, False)
```

```
[ ]: start=int(df_by_month.shape[0]*0.8)
      train=df_by_month[:start]
      test=df_by_month[start:]
      plt.figure(figsize=fig_size)
      plt.plot(train[target])
      plt.plot(test[target])
      plt.title('Training and Testing split - '+ crime)
      plt.show()
```



```
[ ]: model=auto_arima(train[target],start_p=0,d=1,start_q=0,
    max_p=10,max_d=10,max_q=10, start_P=0,
    D=1, start_Q=0, max_P=10,max_D=10,
    max_Q=10, m=12, seasonal=True,
    error_action='warn',trace=True,
    supress_warnings=True,stepwise=True,
    random_state=20,n_fits=50)
```

Performing stepwise search to minimize aic

ARIMA(0,1,0)(0,1,0)[12]	: AIC=1589.575, Time=0.06 sec
ARIMA(1,1,0)(1,1,0)[12]	: AIC=1538.297, Time=0.71 sec
ARIMA(0,1,1)(0,1,1)[12]	: AIC=1493.290, Time=4.14 sec
ARIMA(0,1,1)(0,1,0)[12]	: AIC=1554.976, Time=0.12 sec
ARIMA(0,1,1)(1,1,1)[12]	: AIC=1493.261, Time=3.89 sec
ARIMA(0,1,1)(1,1,0)[12]	: AIC=1528.746, Time=0.56 sec
ARIMA(0,1,1)(2,1,1)[12]	: AIC=1492.553, Time=31.03 sec
ARIMA(0,1,1)(2,1,0)[12]	: AIC=1508.505, Time=5.70 sec
ARIMA(0,1,1)(3,1,1)[12]	: AIC=1493.474, Time=33.00 sec
ARIMA(0,1,1)(2,1,2)[12]	: AIC=1493.588, Time=50.52 sec
ARIMA(0,1,1)(1,1,2)[12]	: AIC=1494.060, Time=28.50 sec
ARIMA(0,1,1)(3,1,0)[12]	: AIC=1495.957, Time=8.14 sec
ARIMA(0,1,1)(3,1,2)[12]	: AIC=1495.349, Time=31.81 sec
ARIMA(0,1,0)(2,1,1)[12]	: AIC=1524.448, Time=3.48 sec
ARIMA(1,1,1)(2,1,1)[12]	: AIC=1490.620, Time=5.36 sec
ARIMA(1,1,1)(1,1,1)[12]	: AIC=1491.348, Time=0.90 sec
ARIMA(1,1,1)(2,1,0)[12]	: AIC=1507.600, Time=1.53 sec
ARIMA(1,1,1)(3,1,1)[12]	: AIC=1491.074, Time=6.20 sec
ARIMA(1,1,1)(2,1,2)[12]	: AIC=1491.691, Time=2.64 sec
ARIMA(1,1,1)(1,1,0)[12]	: AIC=1527.738, Time=0.13 sec
ARIMA(1,1,1)(1,1,2)[12]	: AIC=1492.186, Time=1.65 sec
ARIMA(1,1,1)(3,1,0)[12]	: AIC=1493.405, Time=0.97 sec
ARIMA(1,1,1)(3,1,2)[12]	: AIC=1493.010, Time=3.59 sec
ARIMA(1,1,0)(2,1,1)[12]	: AIC=1503.019, Time=0.85 sec
ARIMA(2,1,1)(2,1,1)[12]	: AIC=1492.105, Time=1.88 sec
ARIMA(1,1,2)(2,1,1)[12]	: AIC=1492.040, Time=2.49 sec
ARIMA(0,1,2)(2,1,1)[12]	: AIC=1491.564, Time=1.11 sec

```

ARIMA(2,1,0)(2,1,1)[12]          : AIC=1497.351, Time=1.12 sec
ARIMA(2,1,2)(2,1,1)[12]          : AIC=1494.499, Time=1.98 sec
ARIMA(1,1,1)(2,1,1)[12] intercept : AIC=1492.326, Time=2.29 sec

```

```

Best model:  ARIMA(1,1,1)(2,1,1)[12]
Total fit time: 236.360 seconds

```

```
[ ]: model.summary()
```

```
[ ]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                                SARIMAX Results
=====
=====
Dep. Variable:                  y    No. Observations:
214
Model:                SARIMAX(1, 1, 1)x(2, 1, 1, 12)    Log Likelihood
-739.310
Date:                  Sun, 23 Apr 2023    AIC
1490.620
Time:                  01:35:32    BIC
1510.440
Sample:                01-31-2001    HQIC
1498.640
                                - 10-31-2018
Covariance Type:                opg
=====

```

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.2987	0.141	2.119	0.034	0.022	0.575
ma.L1	-0.7356	0.098	-7.524	0.000	-0.927	-0.544
ar.S.L12	0.0505	0.126	0.401	0.689	-0.196	0.297
ar.S.L24	-0.1595	0.073	-2.186	0.029	-0.303	-0.016
ma.S.L12	-0.6788	0.104	-6.518	0.000	-0.883	-0.475
sigma2	87.5092	7.234	12.098	0.000	73.332	101.687

```

=====
===
Ljung-Box (L1) (Q):                0.06    Jarque-Bera (JB):
12.46
Prob(Q):                0.81    Prob(JB):
0.00
Heteroskedasticity (H):            0.75    Skew:
0.27
Prob(H) (two-sided):            0.25    Kurtosis:
4.09
=====
===


```


Warnings:

```
[1] Covariance matrix calculated using the outer product of gradients (complex-  
step).
```

```
"""
```

```
[ ]: prediction = pd.DataFrame(model.predict(n_periods = train.shape[0]),index=test.  
    ↪index)  
prediction.columns = ['predicted_crime']  
plt.figure(figsize=fig_size)  
plt.plot(train[target],label="Training")  
plt.plot(test[target],label="Test")  
plt.plot(prediction,label="Predicted")  
plt.legend(loc = 'upper left')  
plt.savefig('../output/%s_%s_pred.jpg' % (city,crime))  
plt.title('Model output - '+crime)  
plt.show()
```



```
[ ]: np.sqrt(np.square(np.subtract(test[target].values,prediction['predicted_crime'].  
    ↪values)).mean())
```

```
[ ]: 45.26006197362022
```