

analysis_merged_type1

April 23, 2023

1 Analysis Template

1.1 Preprocess

```
[ ]: # resolve dependency
     # !pip install pmdarima
```

```
[ ]: import numpy as np
     import pandas as pd
     import matplotlib.pyplot as plt
     from statsmodels.tsa.stattools import adfuller
     from pandas.plotting import autocorrelation_plot
     from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
     import statsmodels.api as sm
     from pmdarima.arima import ADFTest, auto_arima
     %matplotlib inline
```

```
[ ]: data_path = "../data/nypd_assault.csv"
     crime = "type1"
     target = "count"
     date = "date"
     city = "merged"
     fig_size = (20,5)
```

```
[ ]: df_by_day_nyc = pd.read_csv(data_path)
     df_by_day_nyc[date] = pd.to_datetime(df_by_day_nyc[date])
     df_by_day_nyc.set_index(date, inplace=True)
```

```
[ ]: data_path = "../data/battery_occurrence_per_day.csv"
     target = "Count"
     date = "Date"
```

```
[ ]: df_by_day_chi = pd.read_csv(data_path)
     df_by_day_chi[date] = pd.to_datetime(df_by_day_chi[date])
     df_by_day_chi.set_index(date, inplace=True)
```

```
[ ]: df_by_day = df_by_day_nyc.join(df_by_day_chi, how='inner')
```

```
[ ]: df_by_day[target]=df_by_day[target]+df_by_day['count']
df_by_day.drop('count',axis=1,inplace=True)
```

```
[ ]: df_by_day
```

```
[ ]:
      Count
2006-01-01    476
2006-01-02    269
2006-01-03    264
2006-01-04    268
2006-01-05    266
...
2021-12-27    192
2021-12-28    171
2021-12-29    156
2021-12-30    184
2021-12-31    215

[5844 rows x 1 columns]
```

1.2 Profiling

1.2.1 By day


```
[ ]: df_by_day.head()
```

```
[ ]:
      Count
2006-01-01    476
2006-01-02    269
2006-01-03    264
2006-01-04    268
2006-01-05    266
```

```
[ ]: df_by_day.describe()
```

```
[ ]:
      Count
count  5844.000000
mean    289.381588
std      56.037430
min     115.000000
25%     251.000000
50%     286.000000
75%     324.000000
max     531.000000
```

```
[ ]: df_by_day.plot(figsize=fig_size, title="Daily Average Occurrence " + crime)
plt.show()
```




```
[ ]: df_by_day[target].sort_values(ascending=False).head()
```

```
[ ]: 2007-01-01    531
      2011-01-01    498
      2007-07-05    495
      2011-04-10    493
      2013-01-01    481
      Name: Count, dtype: int64
```

1.2.2 By week

```
[ ]: df_by_week = pd.DataFrame(df_by_day[target].resample('W').mean())
```


```
[ ]: df_by_week.plot(
      figsize=fig_size,
      title="Weekly Average Occurrence " + crime)
plt.show()
```



1.2.3 By month

```
[ ]: df_by_month = pd.DataFrame(df_by_day[target].resample('M').mean())
```

```
[ ]: df_by_month.plot(
    figsize=fig_size,
    title="Monthly Average Occurrence " + crime)
plt.show()
```



1.3 Analysis

```
[ ]: #Ho: It is non stationary
    #H1: It is stationary

def adfuller_test(count):
    result=adfuller(count)
    labels = ['ADF Test Statistic','p-value','#Lags Used','Number of_
↳Observations Used']
    for value,label in zip(result,labels):
        print(label+' : '+str(value) )
    if result[1] <= 0.05:
        print("strong evidence against the null hypothesis(Ho), reject the null_
↳hypothesis. Data has no unit root and is stationary")
    else:
        print("weak evidence against null hypothesis, time series has a unit_
↳root, indicating it is non-stationary ")
```

1.3.1 Checking stationary

```
[ ]: adfuller_test(df_by_month[target])
```

```
ADF Test Statistic : -1.1711376179558388
p-value : 0.6858867530399138
#Lags Used : 14
Number of Observations Used : 177
weak evidence against null hypothesis, time series has a unit root, indicating
it is non-stationary
```

1.3.2 Checking seasonality


```
[ ]: df_by_month['seasonal_first_difference'] = df_by_month[target] -  
      ↪df_by_month[target].shift(12)
```

```
[ ]: adfuller_test(df_by_month['seasonal_first_difference'].dropna())
```

```
ADF Test Statistic : -4.057446544331046  
p-value : 0.001137844432949395  
#Lags Used : 12  
Number of Observations Used : 167  
strong evidence against the null hypothesis(Ho), reject the null hypothesis.  
Data has no unit root and is stationary
```

```
[ ]: df_by_month['seasonal_first_difference'].plot(figsize=fig_size, title='Seasonal_  
      ↪difference plot - ' + crime)
```


```
[ ]: <Axes: title={'center': 'Seasonal difference plot - type1'}>
```



1.3.3 Auto Regressive Model

```
[ ]: fig = plt.figure(figsize=(12,8))  
ax1 = fig.add_subplot(211)  
fig = sm.graphics.tsa.plot_acf(df_by_month['seasonal_first_difference'].iloc[13:  
      ↪],lags=80,ax=ax1)  
ax2 = fig.add_subplot(212)  
fig = sm.graphics.tsa.plot_pacf(df_by_month['seasonal_first_difference'].  
      ↪iloc[13:],lags=80,ax=ax2)
```

```
/Users/xuyanchong/opt/anaconda3/lib/python3.9/site-  
packages/statsmodels/graphics/tsaplots.py:348: FutureWarning: The default method  
'yw' can produce PACF values outside of the [-1,1] interval. After 0.13, the  
default will change to unadjusted Yule-Walker ('ywm'). You can use this method  
now by setting method='ywm'.  
warnings.warn(
```




1.3.4 Implementing Seasonal Arima Model

```
[ ]: adf_test=ADFTTest(alpha=0.05)
      adf_test.should_diff(df_by_month[target])
```

```
[ ]: (0.01, False)
```

```
[ ]: start=int(df_by_month.shape[0]*0.8)
      train=df_by_month[:start]
      test=df_by_month[start:]
      plt.figure(figsize=fig_size)
      plt.plot(train[target])
      plt.plot(test[target])
      plt.title('Training and Testing split - '+ crime)
      plt.show()
```



```
[ ]: model=auto_arima(train[target],start_p=0,d=1,start_q=0,
    max_p=10,max_d=10,max_q=10, start_P=0,
    D=1, start_Q=0, max_P=10,max_D=10,
    max_Q=10, m=12, seasonal=True,
    error_action='warn',trace=True,
    supress_warnings=True,stepwise=True,
    random_state=20,n_fits=50)
```

Performing stepwise search to minimize aic

ARIMA(0,1,0)(0,1,0)[12]	: AIC=1163.271, Time=0.14 sec
ARIMA(1,1,0)(1,1,0)[12]	: AIC=1109.836, Time=0.52 sec
ARIMA(0,1,1)(0,1,1)[12]	: AIC=1075.398, Time=2.77 sec
ARIMA(0,1,1)(0,1,0)[12]	: AIC=1110.967, Time=0.22 sec
ARIMA(0,1,1)(1,1,1)[12]	: AIC=1076.900, Time=4.57 sec
ARIMA(0,1,1)(0,1,2)[12]	: AIC=1076.813, Time=14.82 sec
ARIMA(0,1,1)(1,1,0)[12]	: AIC=1089.929, Time=0.97 sec
ARIMA(0,1,1)(1,1,2)[12]	: AIC=inf, Time=50.71 sec
ARIMA(0,1,0)(0,1,1)[12]	: AIC=1115.759, Time=1.24 sec
ARIMA(1,1,1)(0,1,1)[12]	: AIC=1074.612, Time=2.35 sec
ARIMA(1,1,1)(0,1,0)[12]	: AIC=1112.861, Time=0.22 sec
ARIMA(1,1,1)(1,1,1)[12]	: AIC=1076.241, Time=3.11 sec
ARIMA(1,1,1)(0,1,2)[12]	: AIC=1076.192, Time=20.96 sec
ARIMA(1,1,1)(1,1,0)[12]	: AIC=1090.601, Time=1.34 sec
ARIMA(1,1,1)(1,1,2)[12]	: AIC=inf, Time=60.67 sec
ARIMA(1,1,0)(0,1,1)[12]	: AIC=1090.474, Time=2.28 sec
ARIMA(2,1,1)(0,1,1)[12]	: AIC=1074.697, Time=2.52 sec
ARIMA(1,1,2)(0,1,1)[12]	: AIC=1074.666, Time=2.81 sec
ARIMA(0,1,2)(0,1,1)[12]	: AIC=1075.568, Time=1.74 sec
ARIMA(2,1,0)(0,1,1)[12]	: AIC=1085.717, Time=2.21 sec
ARIMA(2,1,2)(0,1,1)[12]	: AIC=1076.542, Time=3.59 sec
ARIMA(1,1,1)(0,1,1)[12] intercept	: AIC=1076.401, Time=1.87 sec

Best model: ARIMA(1,1,1)(0,1,1)[12]

Total fit time: 181.636 seconds

```
[ ]: model.summary()
```

```
[ ]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                                SARIMAX Results
=====
Dep. Variable:                  y    No. Observations:
153
Model:              SARIMAX(1, 1, 1)x(0, 1, 1, 12)    Log Likelihood
-533.306
Date:                  Sun, 23 Apr 2023    AIC
1074.612
Time:                  01:34:33    BIC
1086.379
Sample:                01-31-2006    HQIC
1079.394
                        - 09-30-2018
Covariance Type:                opg
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1           0.2391       0.118       2.029     0.042       0.008       0.470
ma.L1          -0.8508       0.061     -14.033     0.000      -0.970      -0.732
ma.S.L12       -0.6594       0.096      -6.879     0.000      -0.847      -0.472
sigma2         112.6244      14.294       7.879     0.000      84.609     140.640
=====
===
Ljung-Box (L1) (Q):                0.09    Jarque-Bera (JB):
2.51
Prob(Q):                0.76    Prob(JB):
0.28
Heteroskedasticity (H):            0.56    Skew:
-0.33
Prob(H) (two-sided):            0.05    Kurtosis:
2.98
=====
===
```


Warnings:

```
[1] Covariance matrix calculated using the outer product of gradients (complex-
step).
      """
```

```
[ ]: prediction = pd.DataFrame(model.predict(n_periods = train.shape[0]),index=test.
      ↪index)
      prediction.columns = ['predicted_crime']
```



```
plt.figure(figsize=fig_size)
plt.plot(train[target],label="Training")
plt.plot(test[target],label="Test")
plt.plot(prediction,label="Predicted")
plt.legend(loc = 'upper left')
plt.savefig('../output/%s_%s_pred.jpg' % (city,crime))
plt.title('Model output - '+crime)
plt.show()
```



```
[ ]: np.sqrt(np.square(np.subtract(test[target].values,prediction['predicted_crime'].
↪values))).mean())
```

```
[ ]: 38.55733902461584
```