

# analysis\_nyc\_type1

April 23, 2023

## 1 Analysis Template

### 1.1 Preprocess

```
[ ]: # resolve dependency
!pip install pmdarima
```

```
Requirement already satisfied: pmdarima in
/Users/xuyanchong/opt/anaconda3/lib/python3.9/site-packages (2.0.3)
Requirement already satisfied: setuptools!=50.0.0,>=38.6.0 in
/Users/xuyanchong/opt/anaconda3/lib/python3.9/site-packages (from pmdarima)
(65.6.3)
Requirement already satisfied: joblib>=0.11 in
/Users/xuyanchong/opt/anaconda3/lib/python3.9/site-packages (from pmdarima)
(1.1.1)
Requirement already satisfied: statsmodels>=0.13.2 in
/Users/xuyanchong/opt/anaconda3/lib/python3.9/site-packages (from pmdarima)
(0.13.5)
Requirement already satisfied: scikit-learn>=0.22 in
/Users/xuyanchong/opt/anaconda3/lib/python3.9/site-packages (from pmdarima)
(1.0.2)
Requirement already satisfied: Cython!=0.29.18,!0.29.31,>=0.29 in
/Users/xuyanchong/opt/anaconda3/lib/python3.9/site-packages (from pmdarima)
(0.29.33)
Requirement already satisfied: urllib3 in
/Users/xuyanchong/opt/anaconda3/lib/python3.9/site-packages (from pmdarima)
(1.26.14)
Requirement already satisfied: numpy>=1.21.2 in
/Users/xuyanchong/opt/anaconda3/lib/python3.9/site-packages (from pmdarima)
(1.22.4)
Requirement already satisfied: pandas>=0.19 in
/Users/xuyanchong/opt/anaconda3/lib/python3.9/site-packages (from pmdarima)
(1.5.3)
Requirement already satisfied: scipy>=1.3.2 in
/Users/xuyanchong/opt/anaconda3/lib/python3.9/site-packages (from pmdarima)
(1.7.1)
Requirement already satisfied: python-dateutil>=2.8.1 in
/Users/xuyanchong/opt/anaconda3/lib/python3.9/site-packages (from
```

```

pandas>=0.19->pmdarima) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in
/Users/xuyanchong/opt/anaconda3/lib/python3.9/site-packages (from
pandas>=0.19->pmdarima) (2022.7)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/Users/xuyanchong/opt/anaconda3/lib/python3.9/site-packages (from scikit-
learn>=0.22->pmdarima) (2.2.0)
Requirement already satisfied: packaging>=21.3 in
/Users/xuyanchong/opt/anaconda3/lib/python3.9/site-packages (from
statsmodels>=0.13.2->pmdarima) (23.0)
Requirement already satisfied: patsy>=0.5.2 in
/Users/xuyanchong/opt/anaconda3/lib/python3.9/site-packages (from
statsmodels>=0.13.2->pmdarima) (0.5.3)
Requirement already satisfied: six in
/Users/xuyanchong/opt/anaconda3/lib/python3.9/site-packages (from
patsy>=0.5.2->statsmodels>=0.13.2->pmdarima) (1.16.0)

```

```

[ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from statsmodels.tsa.stattools import adfuller
from pandas.plotting import autocorrelation_plot
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
import statsmodels.api as sm
from pmdarima.arima import ADFTest , auto_arima
%matplotlib inline

```

```

[ ]: data_path = "../data/nypd_assault.csv"
crime = "type1"
target = "count"
date = "date"
city = "nyc"
fig_size = (20,5)

```

```

[ ]: df_by_day = pd.read_csv(data_path)
df_by_day[date] = pd.to_datetime(df_by_day[date])
df_by_day.set_index(date, inplace=True)

```

## 1.2 Profiling

### 1.2.1 By day

```

[ ]: df_by_day.head()

```

```

[ ]:
      count
date
2006-01-01    107


```

```
2006-01-02    105
2006-01-03    104
2006-01-04    118
2006-01-05    133
```

```
[ ]: df_by_day.describe()
```

```
[ ]:
      count
count  5844.000000
mean   131.463895
std    24.073300
min     32.000000
25%    116.000000
50%    132.000000
75%    148.000000
max    221.000000
```

```
[ ]: df_by_day.plot(figsize=fig_size, title="Daily Average Occurrence " + crime)
plt.show()
```




```
[ ]: df_by_day[target].sort_values(ascending=False).head()
```

```
[ ]: date
2014-06-11    221
2014-01-01    216
2011-01-01    213
2016-01-01    212
2014-06-04    211
Name: count, dtype: int64
```

### 1.2.2 By week

```
[ ]: df_by_week = pd.DataFrame(df_by_day[target].resample('W').mean())
```


```
[ ]: df_by_week.plot(
    figsize=fig_size,
    title="Weekly Average Occurrence " + crime)
plt.show()
```



### 1.2.3 By month

```
[ ]: df_by_month = pd.DataFrame(df_by_day[target].resample('M').mean())
```

```
[ ]: df_by_month.plot(
    figsize=fig_size,
    title="Monthly Average Occurrence " + crime)
plt.show()
```



## 1.3 Analysis

```
[ ]: #Ho: It is non stationary
    #H1: It is stationary

def adfuller_test(count):
    result=adfuller(count)
```

```

labels = ['ADF Test Statistic','p-value','#Lags Used','Number of_
↳Observations Used']
for value,label in zip(result,labels):
    print(label+' : '+str(value) )
if result[1] <= 0.05:
    print("strong evidence against the null hypothesis(Ho), reject the null_
↳hypothesis. Data has no unit root and is stationary")
else:
    print("weak evidence against null hypothesis, time series has a unit_
↳root, indicating it is non-stationary ")

```

### 1.3.1 Checking stationary

```
[ ]: adfuller_test(df_by_month[target])
```

```

ADF Test Statistic : -1.8558505343444571
p-value : 0.3531067118505653
#Lags Used : 14
Number of Observations Used : 177
weak evidence against null hypothesis, time series has a unit root, indicating
it is non-stationary

```

### 1.3.2 Checking seasonality

```
[ ]: df_by_month['seasonal_first_difference'] = df_by_month[target] -_
↳df_by_month[target].shift(12)
```

```
[ ]: adfuller_test(df_by_month['seasonal_first_difference'].dropna())
```


```

ADF Test Statistic : -3.516023909094369
p-value : 0.007584977307881228
#Lags Used : 12
Number of Observations Used : 167
strong evidence against the null hypothesis(Ho), reject the null hypothesis.
Data has no unit root and is stationary

```

```
[ ]: df_by_month['seasonal_first_difference'].plot(figsize=fig_size, title='Seasonal_
↳difference plot - ' + crime)
```


```
[ ]: <Axes: title={'center': 'Seasonal difference plot - type1'}, xlabel='date'>
```



### 1.3.3 Auto Regressive Model

```
[ ]: fig = plt.figure(figsize=(12,8))
      ax1 = fig.add_subplot(211)
      fig = sm.graphics.tsa.plot_acf(df_by_month['seasonal_first_difference'].iloc[13:
      ↪],lags=80,ax=ax1)
      ax2 = fig.add_subplot(212)
      fig = sm.graphics.tsa.plot_pacf(df_by_month['seasonal_first_difference'].
      ↪iloc[13:],lags=80,ax=ax2)
```

/Users/xuyanchong/opt/anaconda3/lib/python3.9/site-packages/statsmodels/graphics/tsaplots.py:348: FutureWarning: The default method 'yw' can produce PACF values outside of the [-1,1] interval. After 0.13, the default will change to unadjusted Yule-Walker ('ywm'). You can use this method now by setting method='ywm'.  
 warnings.warn(




### 1.3.4 Implementing Seasonal Arima Model

```
[ ]: adf_test=ADFTTest(alpha=0.05)
      adf_test.should_diff(df_by_month[target])
```

```
[ ]: (0.01, False)
```

```
[ ]: start=int(df_by_month.shape[0]*0.8)
      train=df_by_month[:start]
      test=df_by_month[start:]
      plt.figure(figsize=fig_size)
      plt.plot(train[target])
      plt.plot(test[target])
      plt.title('Training and Testing split - '+ crime)
      plt.show()
```



```
[ ]: model=auto_arima(train[target],start_p=0,d=1,start_q=0,
    max_p=10,max_d=10,max_q=10, start_P=0,
    D=1, start_Q=0, max_P=10,max_D=10,
    max_Q=10, m=12, seasonal=True,
    error_action='warn',trace=True,
    supress_warnings=True,stepwise=True,
    random_state=20,n_fits=50)
```

Performing stepwise search to minimize aic

```
ARIMA(0,1,0)(0,1,0)[12] : AIC=1001.756, Time=0.02 sec
ARIMA(1,1,0)(1,1,0)[12] : AIC=939.636, Time=0.16 sec
ARIMA(0,1,1)(0,1,1)[12] : AIC=895.705, Time=0.33 sec
ARIMA(0,1,1)(0,1,0)[12] : AIC=956.026, Time=0.05 sec
ARIMA(0,1,1)(1,1,1)[12] : AIC=897.560, Time=0.59 sec
ARIMA(0,1,1)(0,1,2)[12] : AIC=897.538, Time=2.41 sec
ARIMA(0,1,1)(1,1,0)[12] : AIC=922.332, Time=0.14 sec
ARIMA(0,1,1)(1,1,2)[12] : AIC=899.294, Time=8.19 sec
ARIMA(0,1,0)(0,1,1)[12] : AIC=928.833, Time=0.32 sec
ARIMA(1,1,1)(0,1,1)[12] : AIC=894.280, Time=0.93 sec
ARIMA(1,1,1)(0,1,0)[12] : AIC=956.773, Time=0.19 sec
ARIMA(1,1,1)(1,1,1)[12] : AIC=895.795, Time=1.74 sec
ARIMA(1,1,1)(0,1,2)[12] : AIC=895.739, Time=21.69 sec
ARIMA(1,1,1)(1,1,0)[12] : AIC=920.270, Time=1.98 sec
ARIMA(1,1,1)(1,1,2)[12] : AIC=inf, Time=62.84 sec
ARIMA(1,1,0)(0,1,1)[12] : AIC=906.736, Time=2.30 sec
ARIMA(2,1,1)(0,1,1)[12] : AIC=894.901, Time=3.52 sec
ARIMA(1,1,2)(0,1,1)[12] : AIC=894.978, Time=3.57 sec
ARIMA(0,1,2)(0,1,1)[12] : AIC=895.509, Time=3.01 sec
ARIMA(2,1,0)(0,1,1)[12] : AIC=903.013, Time=2.43 sec
ARIMA(2,1,2)(0,1,1)[12] : AIC=896.866, Time=4.01 sec
ARIMA(1,1,1)(0,1,1)[12] intercept : AIC=894.087, Time=2.32 sec
ARIMA(1,1,1)(0,1,0)[12] intercept : AIC=958.715, Time=0.09 sec
ARIMA(1,1,1)(1,1,1)[12] intercept : AIC=895.619, Time=2.64 sec
ARIMA(1,1,1)(0,1,2)[12] intercept : AIC=895.561, Time=28.71 sec
ARIMA(1,1,1)(1,1,0)[12] intercept : AIC=922.041, Time=2.32 sec
ARIMA(1,1,1)(1,1,2)[12] intercept : AIC=897.077, Time=54.58 sec
```



```

ARIMA(0,1,1)(0,1,1)[12] intercept : AIC=897.075, Time=0.67 sec
ARIMA(1,1,0)(0,1,1)[12] intercept : AIC=908.535, Time=0.61 sec
ARIMA(2,1,1)(0,1,1)[12] intercept : AIC=inf, Time=1.55 sec
ARIMA(1,1,2)(0,1,1)[12] intercept : AIC=inf, Time=1.56 sec
ARIMA(0,1,0)(0,1,1)[12] intercept : AIC=930.742, Time=0.46 sec
ARIMA(0,1,2)(0,1,1)[12] intercept : AIC=896.221, Time=0.78 sec
ARIMA(2,1,0)(0,1,1)[12] intercept : AIC=904.748, Time=0.72 sec
ARIMA(2,1,2)(0,1,1)[12] intercept : AIC=894.755, Time=1.49 sec

```

Best model: ARIMA(1,1,1)(0,1,1)[12] intercept  
Total fit time: 218.919 seconds

```
[ ]: model.summary()
```

```
[ ]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

#### SARIMAX Results

```

=====
=====
Dep. Variable:                y    No. Observations:
153
Model:                SARIMAX(1, 1, 1)x(0, 1, 1, 12)    Log Likelihood
-442.043
Date:                Sun, 23 Apr 2023    AIC
894.087
Time:                01:34:52    BIC
908.795
Sample:                01-31-2006    HQIC
900.064
                        - 09-30-2018
Covariance Type:                opg
=====

```

	coef	std err	z	P> z	[0.025	0.975]
intercept	-0.0330	0.020	-1.658	0.097	-0.072	0.006
ar.L1	0.3398	0.095	3.586	0.000	0.154	0.526
ma.L1	-0.8941	0.056	-15.964	0.000	-1.004	-0.784
ma.S.L12	-0.8122	0.085	-9.562	0.000	-0.979	-0.646
sigma2	29.2120	3.730	7.831	0.000	21.901	36.523

```

=====
===
Ljung-Box (L1) (Q):                0.21    Jarque-Bera (JB):
20.56
Prob(Q):                0.65    Prob(JB):
0.00
Heteroskedasticity (H):            1.55    Skew:
-0.59

```


Prob(H) (two-sided): 0.14 Kurtosis:  
4.47

=====  
===

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).  
"""

```
[ ]: prediction = pd.DataFrame(model.predict(n_periods = train.shape[0]),index=test.  
    ↪index)  
prediction.columns = ['predicted_crime']  
plt.figure(figsize=fig_size)  
plt.plot(train[target],label="Training")  
plt.plot(test[target],label="Test")  
plt.plot(prediction,label="Predicted")  
plt.legend(loc = 'upper left')  
plt.savefig('../output/%s_%s_pred.jpg' % (city,crime))  
plt.title('Model output - '+crime)  
plt.show()
```



```
[ ]: np.sqrt(np.square(np.subtract(test[target].values,prediction['predicted_crime'].  
    ↪values)).mean())
```

```
[ ]: 19.612174106151354
```